

Lecture2 – Relational Models and Database Schema

ICDT 1202Y – Database Systems

LEARNING OUTCOME

After this session you will be able to :

- Differentiate between the various data models
- Explain the properties of a relation
- Define relational keys
- Explain relational constraints
- Differentiate between database instances and schemas
- Describe a two tier and three tier architecture



DATA MODELS

- A ***data model*** is a collection of high level description concepts for describing data. It hides many low-level storage details
- It is a set of concepts to describe the *structure* of a database, and certain *constraints* that the database should obey.
- A ***schema*** is a description of a particular collection of data, using the given data model

CATEGORIES OF DATA MODELS

- **Conceptual (high-level, semantic) data models:**

- Provides concepts that are close to the way many users perceive data.
- Uses concepts such as *entities*, *attributes* and *relationships*.
- An **entity** represents a real-world object or concept e.g student or lecturer
- An **attribute** represents some property of interest that further describes an entity.
- A **relationship** among two entities represents an association between two entities e.g. teaches relationship.

CATEGORIES OF DATA MODELS

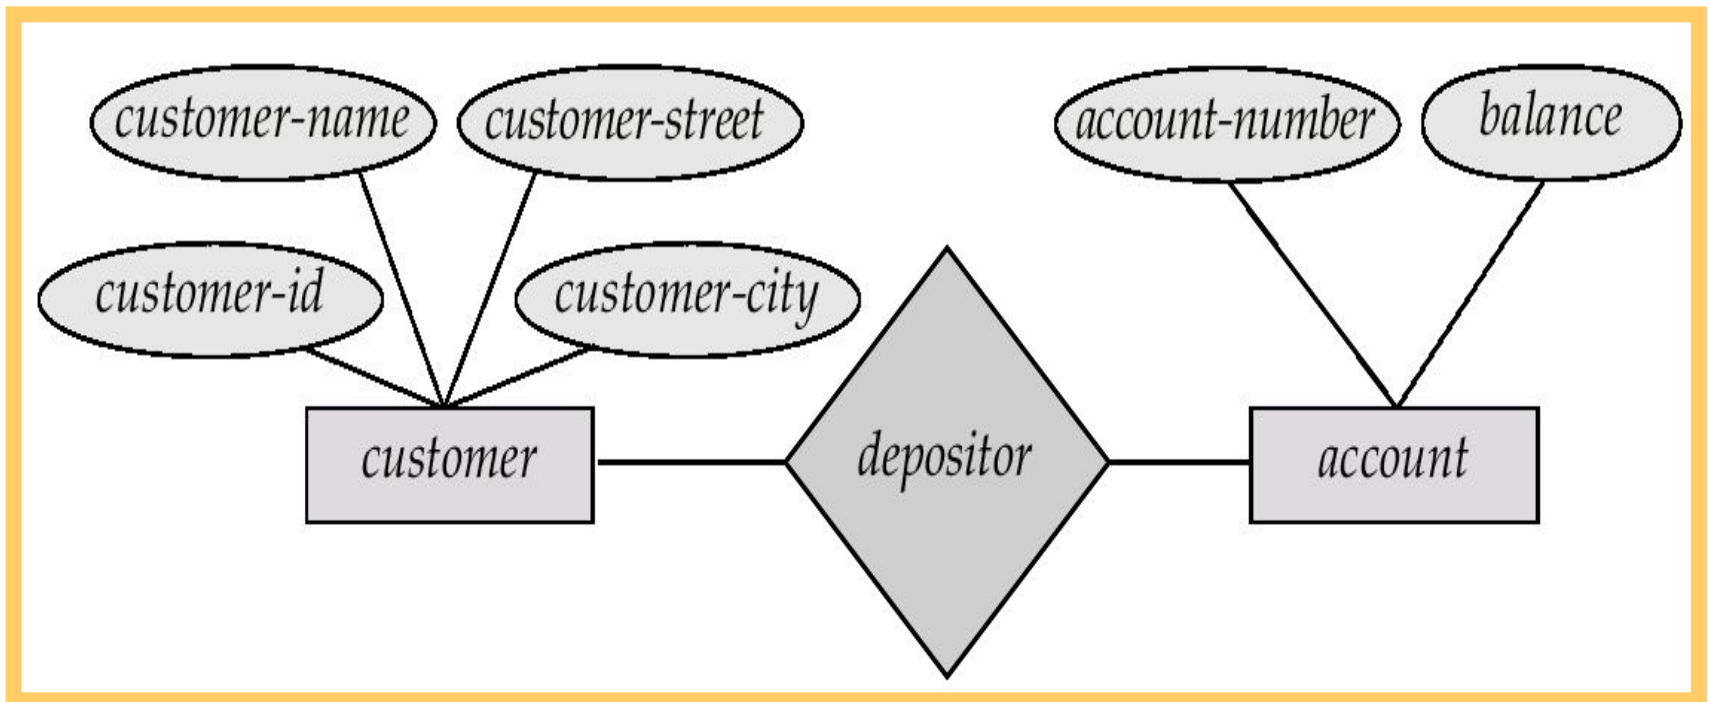
- **Physical (low-level, internal)** data models:
 - Provide concepts that describe details of how data is stored in the computer.
- **Implementation (representational)** data models:
 - Provide concepts that fall between the above two, balancing user views with some computer storage details.
 - These include the widely-used relational data model and network and hierarchical data models.

ENTITY RELATIONSHIP MODEL

- E-R model of real world
- Entities (objects)
 - E.g. customers, accounts, bank branch
- Relationships between entities
 - E.g. Account A-101 is held by customer Johnson
 - Relationship set *depositor* associates customers with accounts
- Widely used for database design
- Database design in E-R model usually converted to design in the relational model (coming up next) which is used for storage and processing

ENTITY RELATIONSHIP MODEL

- Example of schema in the entity- relationship model



RELATIONAL MODEL

- A simple data model: the Relational Data Model - data stored in relations (tables)
- Schema: *RelationName(field1 : type1 , ..., fieldn : typen)*
- A declarative query language: SQL
 - SELECT balance
 - FROM account
 - WHERE branch = 'Springfield'
- Programmer specified what answers a query should return, but not how the query is executed
- DBMS picks the best execution strategy
- Provides physical data independence (applications need not need to worry about how data is physically structured and stored)

RELATIONAL MODEL

- The relational database model achieves structural independence - any type of association be it one-to-one, one-to-many, many-to-many can be easily implemented with the relational model.
- The relational database model has a very powerful and flexible query capability.

RELATIONAL MODEL – EXAMPLE 2

- A relation of students:
Students(sid : string, name : string, age : integer, gpa : real)
- An instance of the students relation can be represented as follows:

<i>sid</i>	<i>name</i>	<i>login</i>	<i>age</i>	<i>gpa</i>
53666	Jones	Jonescs	18	7.4
53668	Smith	smithee	18	7.8
53650	Smith	smithmath	19	6.4
53831	Madayan	madayan@music	11	8.0
53832	Guldy	guldu@music	12	2.0

RELATIONAL MODEL – EXAMPLE 2

- A sample relational database:

<i>customer-id</i>	<i>customer-name</i>	<i>customer-street</i>	<i>customer-city</i>
192-83-7465	Johnson	12 Alma St.	Palo Alto
019-28-3746	Smith	4 North St.	Rye
677-89-9011	Hayes	3 Main St.	Harrison
182-73-6091	Turner	123 Putnam Ave.	Stamford
321-12-3123	Jones	100 Main St.	Harrison
336-66-9999	Lindsay	175 Park Ave.	Pittsfield
019-28-3746	Smith	72 North St.	Rye

(a) The *customer* table

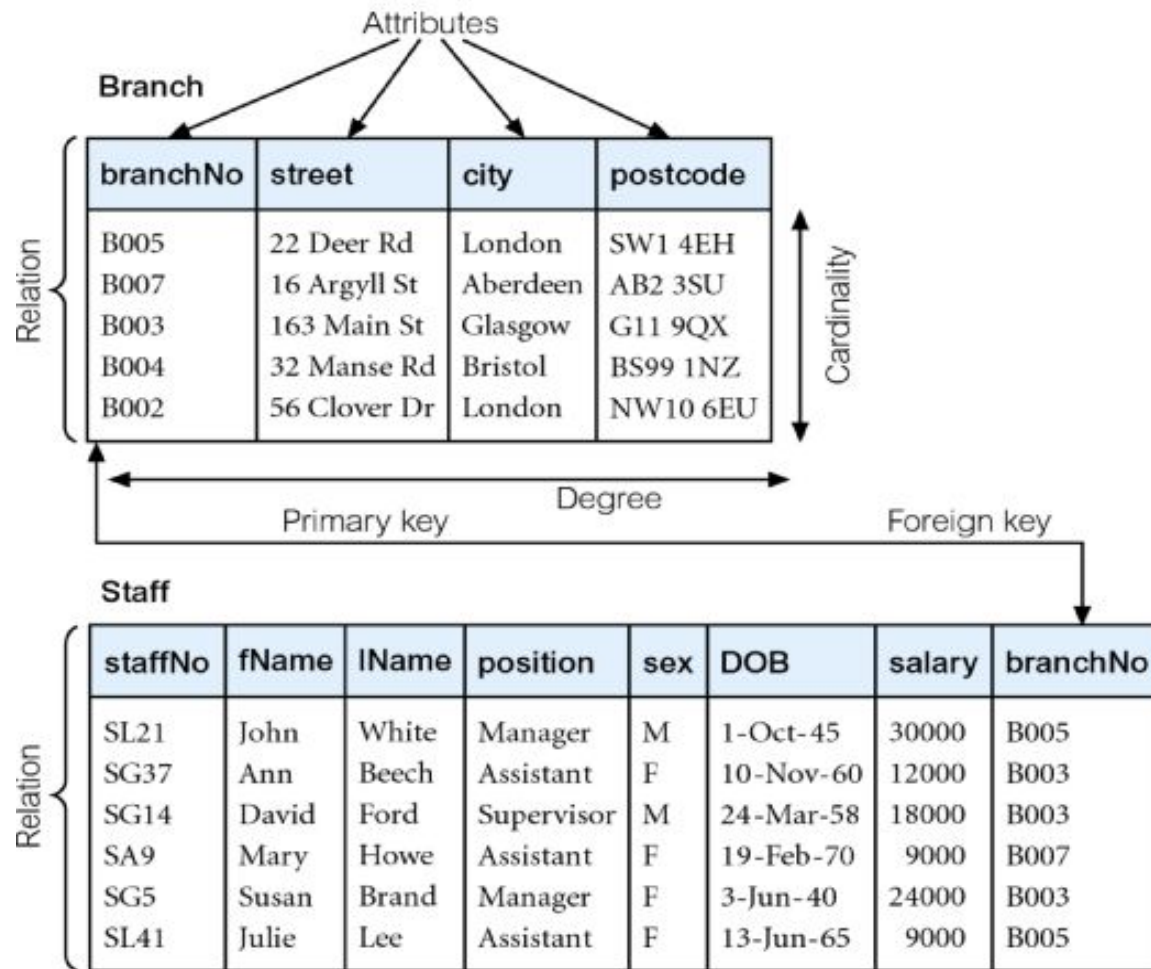
<i>account-number</i>	<i>balance</i>
A-101	500
A-215	700
A-102	400
A-305	350
A-201	900
A-217	750
A-222	700

(b) The *account* table

<i>customer-id</i>	<i>account-number</i>
192-83-7465	A-101
192-83-7465	A-201
019-28-3746	A-215
677-89-9011	A-102
182-73-6091	A-305
321-12-3123	A-217
336-66-9999	A-222
019-28-3746	A-201

(c) The *depositor* table

EXAMPLE RELATION



EXAMPLE OF ATTRIBUTE DOMAINS

Attribute	Domain Name	Meaning	Domain Definition
branchNo	BranchNumbers	The set of all possible branch numbers	character: size 4, range B001–B999
street	StreetNames	The set of all street names in Britain	character: size 25
city	CityNames	The set of all city names in Britain	character: size 15
postcode	Postcodes	The set of all postcodes in Britain	character: size 8
sex	Sex	The sex of a person	character: size 1, value M or F
DOB	DatesOfBirth	Possible values of staff birth dates	date, range from 1-Jan-20, format dd-mmm-yy
salary	Salaries	Possible values of staff salaries	monetary: 7 digits, range 6000.00–40000.00

DEFINITIONS

<u>Informal Terms</u>		<u>Formal Terms</u>
Table		Relation
Column		Attribute/Domain
Row		Tuple
Values in a column		Domain
Table Definition		Schema of a Relation
Populated Table		Extension

RELATIONAL MODEL TERMINOLOGY

- A **relation** is a table with columns and rows.
 - Only applies to logical structure of the database, not the physical structure.
- An **attribute** is a named column of a relation.
- **Domain** is the set of allowable values for one or more attributes.
- **Tuple** is a row of a relation.
- **Degree** is the number of attributes in a relation.
- **Cardinality** is the number of tuples in a relation.
- **Relational Database** is a collection of normalized relations with distinct relation names.

PROPERTIES OF RELATIONS

- Relation name is distinct from all other relation names in relational schema.
- Each cell of relation contains exactly one atomic (single) value.
- Each attribute has a distinct name.
- Values of an attribute are all from the same domain.
- Each tuple is distinct; there are no duplicate tuples.
- Order of attributes has no significance.
- Order of tuples has no significance, theoretically.

RELATIONAL KEYS

- Superkey
 - An attribute, or set of attributes, that uniquely identifies a tuple within a relation.
- Candidate Key
 - Superkey (K) such that no proper subset is a superkey within the relation.
- Primary Key
 - Candidate key selected to identify tuples uniquely within relation.
- Alternate Keys
 - Candidate keys that are not selected to be primary key.
- Foreign Key
 - Attribute, or set of attributes, within one relation that matches candidate key of some (possibly same) relation.

REFERENTIAL INTEGRITY

- A constraint involving *two* relations (the previous constraints involve a *single* relation).
- Used to specify a *relationship* among tuples in two relations: the **referencing relation** and the **referenced relation**.
- Tuples in the *referencing relation* R_1 have attributes FK (called **foreign key** attributes) that reference the primary key attributes PK of the *referenced relation* R_2 . A tuple t_1 in R_1 is said to **reference** a tuple t_2 in R_2 if $t_1[\text{FK}] = t_2[\text{PK}]$.
- A referential integrity constraint can be displayed in a relational database schema as a directed arc from $R_1.\text{FK}$ to $R_2.\text{PK}$

RELATIONAL INTEGRITY CONSTRAINTS

Statement of the constraint

The value in the foreign key column (or columns) FK of the **referencing relation** R_1 can be either:

- (1) a value of an existing primary key value of the corresponding primary key PK in the **referenced relation** R_2 , or..
- (2) a null.

In case (2), the FK in R_1 should not be a part of its own primary key.

SEMANTIC INTEGRITY CONSTRAINTS

- Based on application semantics and cannot be expressed by the model per se
- E.g., “the max. no. of hours per employee for all projects he or she works on is 56 hrs per week”
- A *constraint specification language* may have to be used to express these

EXAMPLE SCHEMA

Figure 7.5 Schema diagram for the COMPANY relational database schema; the primary keys are underlined.

EMPLOYEE

FNAME	MINIT	LNAME	<u>SSN</u>	BDATE	ADDRESS	SEX	SALARY	SUPERSSN	DNO
-------	-------	-------	------------	-------	---------	-----	--------	----------	-----

DEPARTMENT

DNAME	<u>DNUMBER</u>	MGRSSN	MGRSTARTDATE
-------	----------------	--------	--------------

DEPT_LOCATIONS

<u>DNUMBER</u>	<u>DLOCATION</u>
----------------	------------------

PROJECT

PNAME	<u>PNUMBER</u>	PLOCATION	DNUM
-------	----------------	-----------	------

WORKS_ON

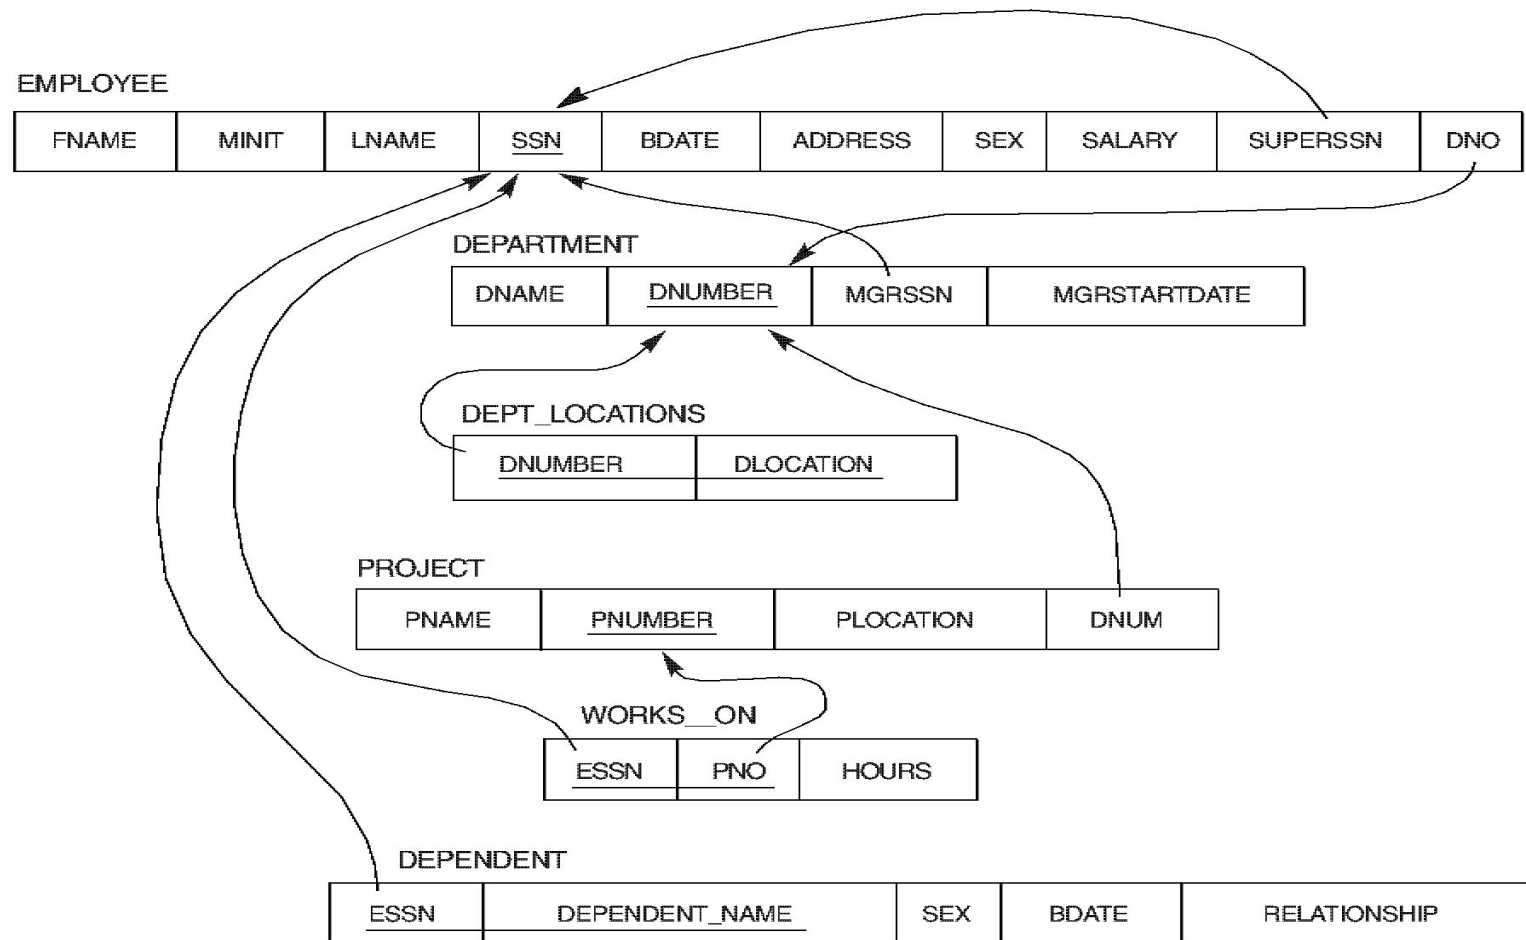
<u>ESSN</u>	<u>PNO</u>	HOURS
-------------	------------	-------

DEPENDENT

<u>ESSN</u>	<u>DEPENDENT_NAME</u>	SEX	BDATE	RELATIONSHIP
-------------	-----------------------	-----	-------	--------------

RELATIONAL INTEGRITY CONSTRAINTS

Figure 7.7 Referential integrity constraints displayed on the COMPANY relational database schema diagram.



Update Operations on Relations

- Update Operations
 - INSERT a tuple.
 - DELETE a tuple.
 - MODIFY a tuple.
- Integrity constraints should not be violated by the update operations.
- Several update operations may have to be grouped together.
- Updates may *propagate* to cause other updates automatically. This may be necessary to maintain integrity constraints.

UPDATE OPERATION ON RELATIONS

- In case of integrity violation, several actions can be taken:
 - Cancel the operation that causes the violation (REJECT option)
 - Perform the operation but inform the user of the violation
 - Trigger additional updates so the violation is corrected (CASCADE option, SET NULL option)
 - Execute a user-specified error-correction routine

EXERCISE

- Consider the following relations for a database that keeps track of student enrollment in courses and the books adopted for each course:
- STUDENT(SSN, Name, Major, Bdate)
- COURSE(Course#, Cname, Dept)
- ENROLL(SSN, Course#, Quarter, Grade)
- BOOK_ADOPTION(Course#, Quarter, Book_ISBN)
- TEXT(Book ISBN, Book_Title, Publisher, Author)
- **Draw a relational schema diagram specifying the foreign keys for this schema.**

DATABASE SCHEMA

- **Database Schema:** The *description* of a database. Includes descriptions of the database structure and the constraints that should hold on the database.
 - Can be represented using a diagram and that is known as a *schema diagram*.
- **Schema Diagram:** A diagrammatic display of (some aspects of) a database schema:
 - Including, names of record types and data items and some types of constraints

DATABASE SCHEMA

- **Schema Construct:** A component of the schema or an object within the schema, e.g., STUDENT, COURSE.
- An example schema diagram:

Student

Name	StudentID	Course
------	-----------	--------

Module

ModuleName	ModuleID	CreditHours	Department
------------	----------	-------------	------------

Grade_Report

studentID	CourseID	Grade
-----------	----------	-------

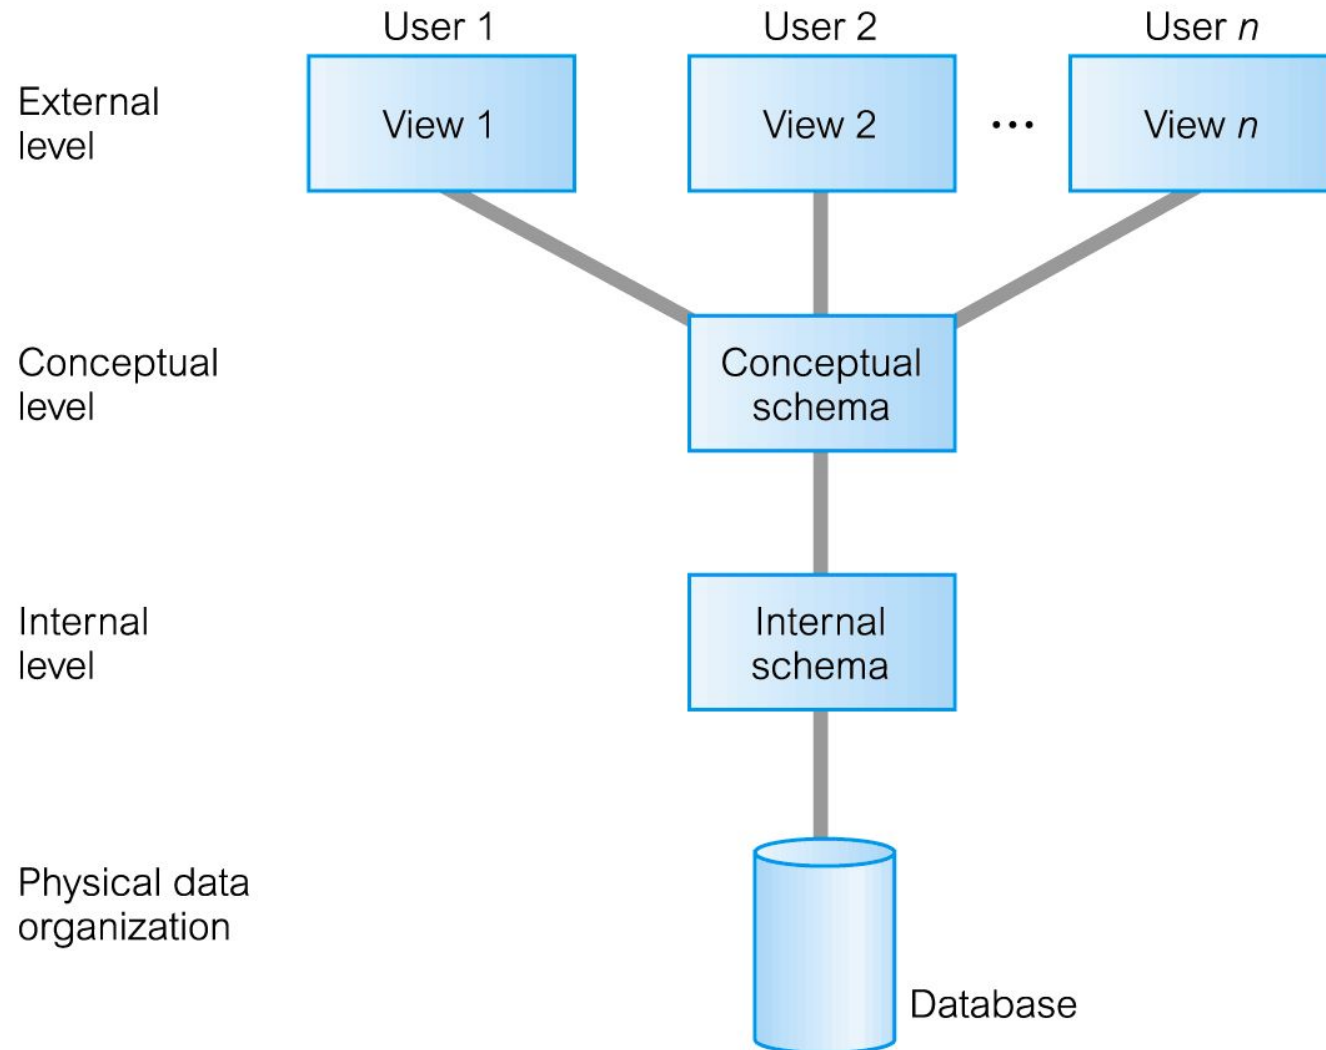
DATABASE INSTANCE/STATE

- **Database Instance/State:** The actual data stored in a database (the content of a database) at a *particular moment in time*.
- For a new database, we define its schema first. At that moment there is no data in the database and it is said to be in an *empty state*.
- **Initial Database State:** Refers to the database when it is loaded with initial data.
- **Valid State:** A state that satisfies the structure and constraints of the database.

DATABASE INSTANCE/STATE

- The **database schema** changes *very infrequently*.
- The **database state** changes *every time the database is updated*.
- **Schema** is also called **intension**, whereas **state** is called **extension**.

THREE-SCHEMA ARCHITECTURE



THREE-SCHEMA ARCHITECTURE

- Many views, single conceptual (logical) schema and physical schema
 - Views (External schema) describe how users see the data
 - Conceptual schema defines logical structure
 - Physical schema describes files and indexes used
- Known as the **Three-Schema Architecture**

THREE-SCHEMA ARCHITECTURE

- Defines DBMS schemas at *three levels*:
 - **Internal schema** at the internal level to describe physical storage structures and access paths. Typically uses a *physical* data model.
 - **Conceptual schema** at the conceptual level to describe the structure and constraints for the *whole* database for a community of users. Uses a *conceptual* or an *implementation* data model.
 - **External schemas** at the external level to describe the various user views. Usually uses the same data model as the conceptual level.

THREE-SCHEMA ARCHITECTURE

External view 1

sNo	fName	lName	age	salary
-----	-------	-------	-----	--------

External view 2

staffNo	lName	branchNo
---------	-------	----------

Conceptual level

staffNo	fName	lName	DOB	salary	branchNo
---------	-------	-------	-----	--------	----------

Internal level

```
struct STAFF {  
    int staffNo;  
    int branchNo;  
    char fName [15];  
    char lName [15];  
    struct date dateOfBirth;  
    float salary;  
    struct STAFF *next;  
};  
index staffNo; index branchNo;
```

/* pointer to next Staff record */
/* define indexes for staff */

CONCEPTUAL SCHEMA

- The conceptual schema describes all the relations stored in the database
- Creating a good conceptual schema is not a simple task. It is called conceptual schema design. It involves:
 - Determining the different relations(entities) needed
 - The number of fields for each relation
 - The type of each field
 - The relationship between relations
 - Constraints
 - ...

INTERNAL SCHEMA

- The internal schema specifies how the relations are actually stored in secondary storage devices
- It also specifies indexes used to speed up access to the relations
- Decisions about the internal schema depend on:
 - Understanding how the data is going to be accessed
 - The facilities provided by the DBMS

EXTERNAL SCHEMA

- The external schema is a refinement of the conceptual schema
- It allows customized and authorized access to individual users or groups of users
- Every database has one physical and one conceptual schema, but *many* external schemas
- Each view is tailored to a particular group of users

DATA INDEPENDENCE

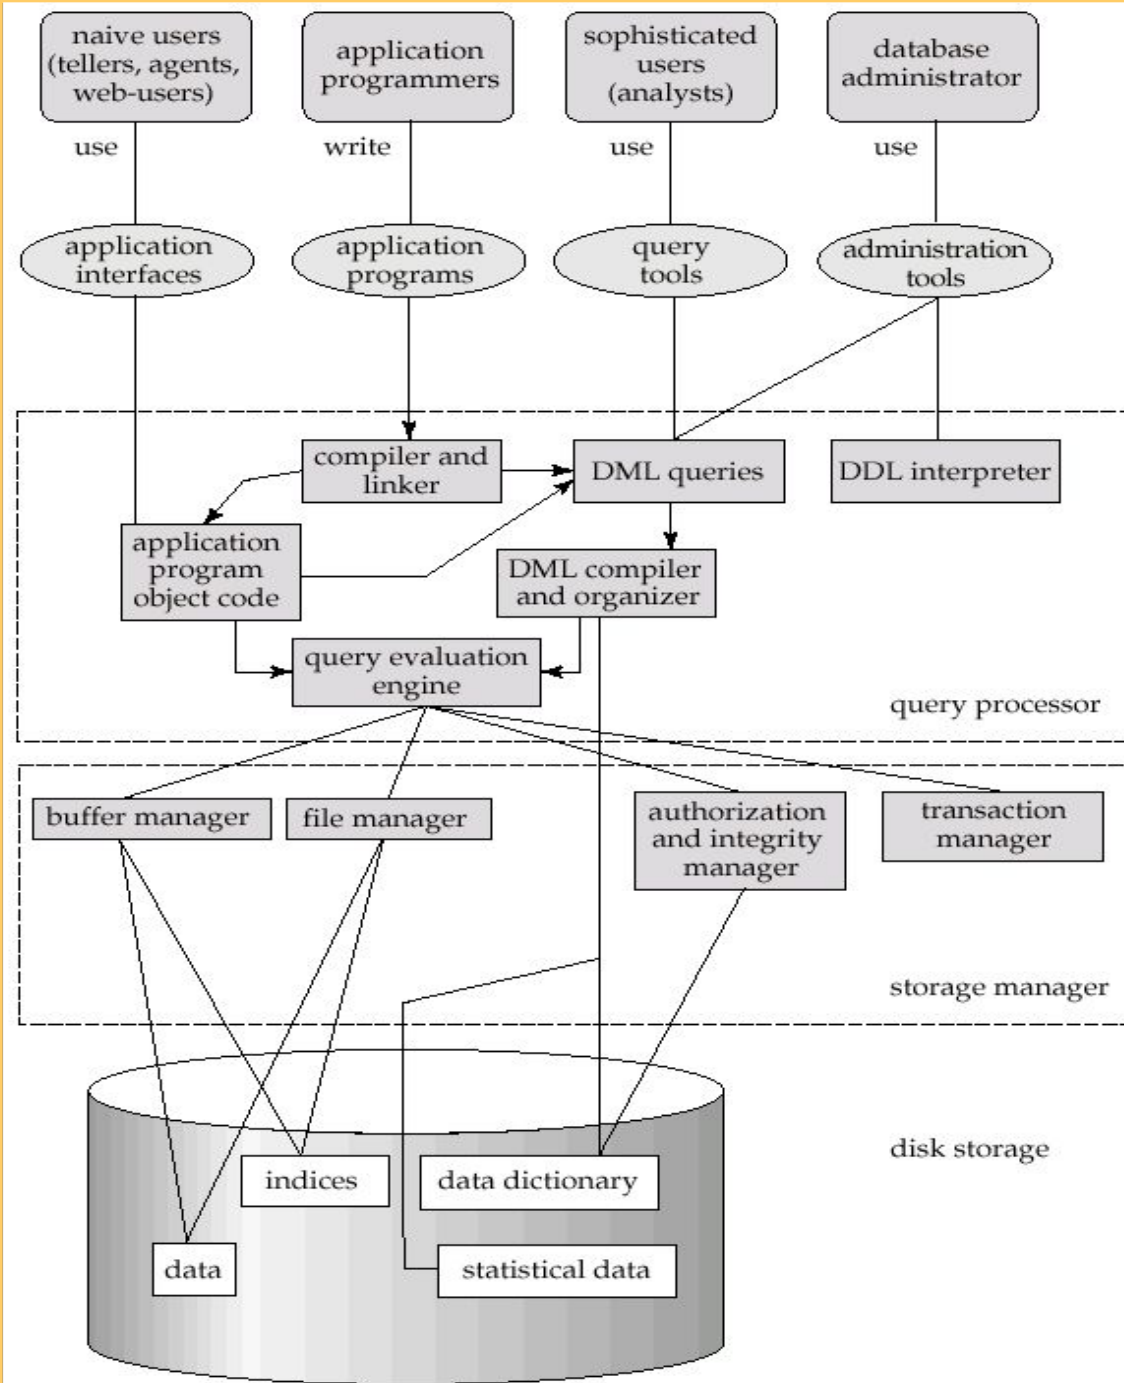
- Applications insulated from how data is structured and stored
- Logical data independence: protection from changes in the *logical* structure of data
 - The capacity to change the conceptual schema without having to change the external schemas and their application programs.
- Physical data independence: protection from changes in the *physical* structure of data
 - The capacity to change the internal schema without having to change the conceptual schema.
- When a schema at a lower level is changed, only the **mappings** between this schema and higher-level schemas need to be changed. The higher-level schemas themselves are *unchanged*.

DBMS LANGUAGES

- **High Level or Non-procedural Languages:** e.g., SQL, are *set-oriented* and specify what data to retrieve than how to retrieve. Also called *declarative* languages.
- **Low Level or Procedural Languages:** record-at-a-time; they specify *how* to retrieve data and include constructs such as looping.

OVERALL SYSTEM ARCHITECTURE

- A database system is partitioned into modules that deal with each of the responsibilities of the system and is broadly divided into **storage manager** and **query processor** components



STORAGE MANAGER

- A program module that provides the interface between the low-level data stored in the db and the app. Programs and queries submitted to the system.
- It is responsible for the interaction with the file manager.
- It translates the various DML statements into low-level file commands.
- It is responsible for the storing, retrieving, and updating data in the db.

STORAGE MANAGER

- It implements several data structures as part of the physical system implementation.
 - **Data files** which store the data itself.
 - **Data dictionary** which stores the metadata about the structure of the db.
 - **Indices** which provide fast access to data items that hold particular values.

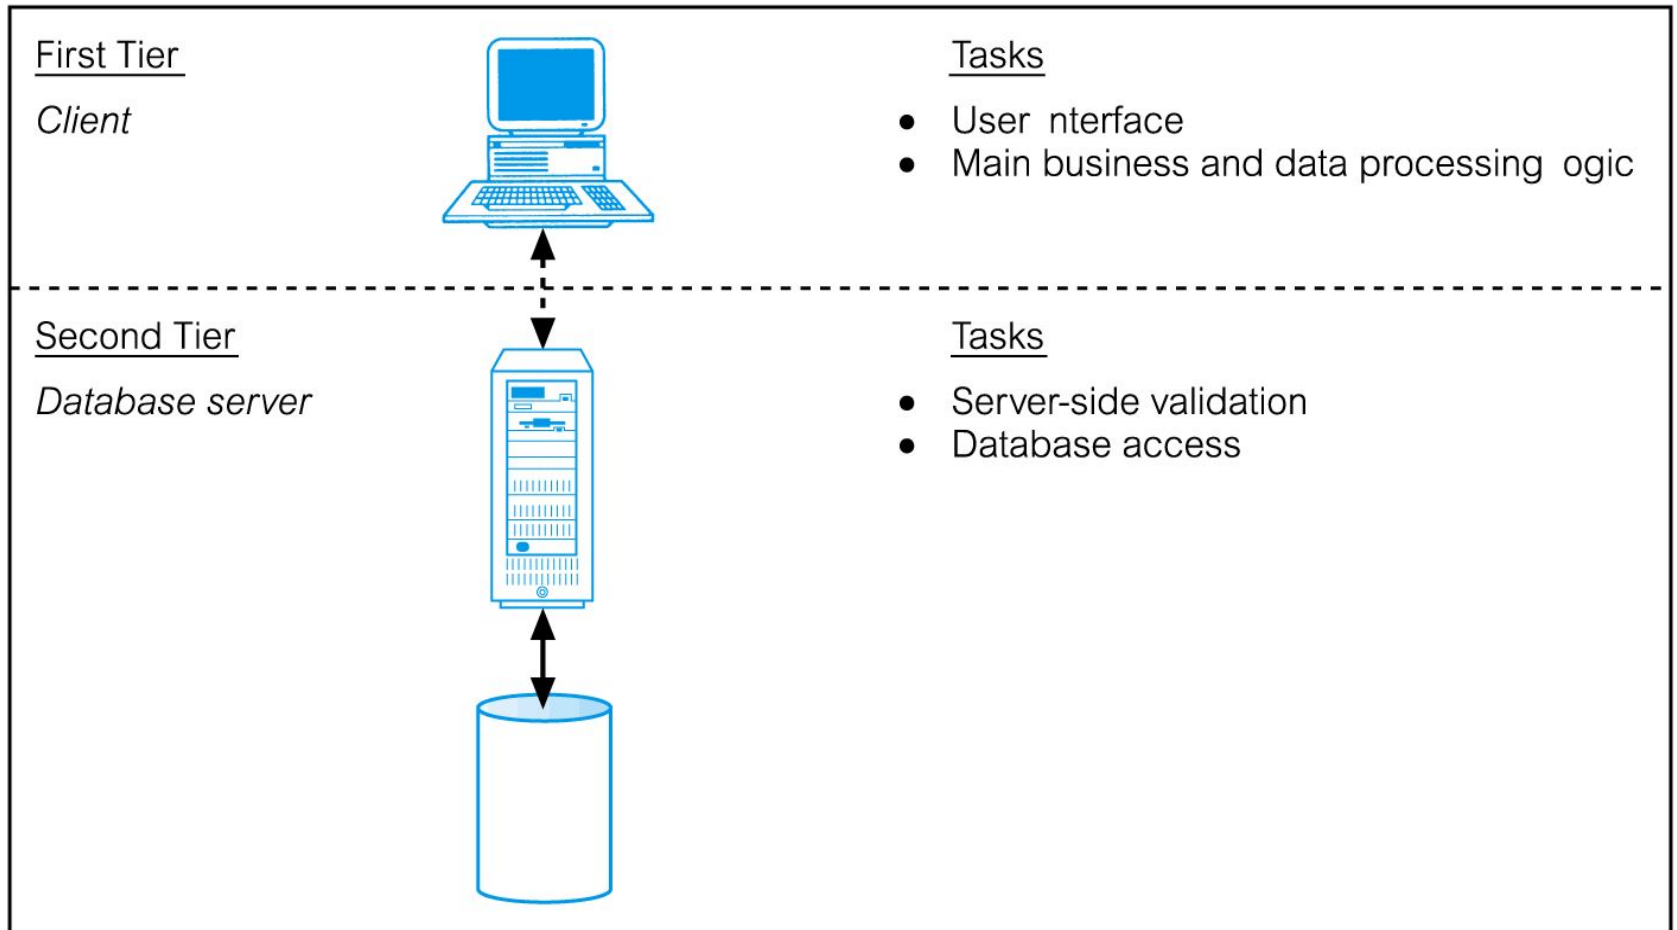
QUERY PROCESSOR

- **DDL interpreter** that interprets DDL statements and records into the definitions in the data dictionary.
- **DML compiler** that translates DML statements in a query language into an evaluation plan consisting of low-level instructions that query evaluation engine understands.
- **Query evaluation engine** which executes low-level instructions generated by the DML compiler.

TWO TIER CLIENT ARCHITECTURE

- Client (tier 1) manages user interface and runs applications.
- Server (tier 2) holds database and DBMS.
- Advantages include:
 - wider access to existing databases;
 - increased performance;
 - possible reduction in hardware costs;
 - reduction in communication costs;
 - increased consistency.

TWO TIER CLIENT ARCHITECTURE



THREE TIER CLIENT ARCHITECTURE

- Client side presented two problems preventing true scalability:
 - Considerable resources required on client's computer to run effectively.
 - Significant client side administration overhead.
- By 1995, three layers proposed, each potentially running on a different platform.
 - Advantages:
 - Client requiring less expensive hardware.
 - Application maintenance centralized.
 - Easier to modify or replace one tier without affecting others.
 - Separating business logic from database functions makes it easier to implement load balancing.
 - Maps quite naturally to Web environment.

THREE TIER CLIENT ARCHITECTURE

